

CIS531 – Compiler Construction

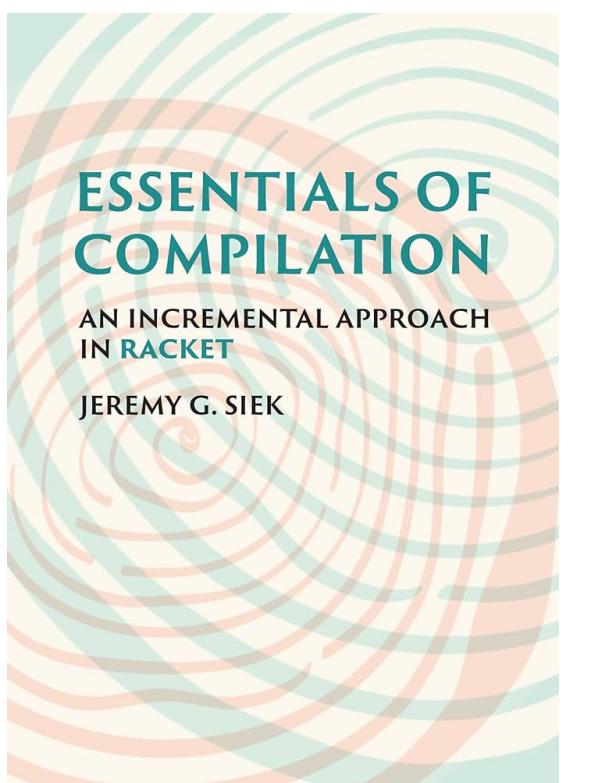
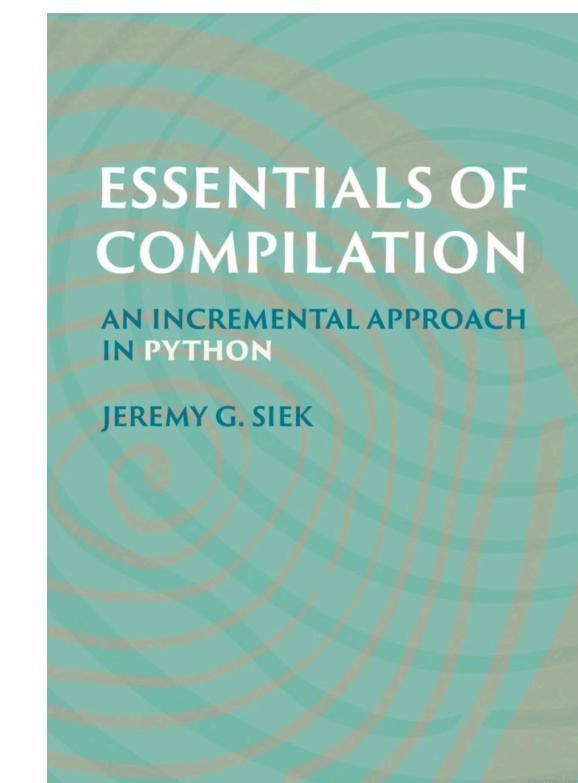
Fall 2025

Kristopher Micinski



Welcome to CIS531!

- Very happy to have you in class! It will be a great semester!
- This course: a tour of compiler design, along with some advanced (not hard!) topics
- We will be following a **free** book, which you should download:
 - There are **several versions of the book**
 - *Project starters were rewritten (sometimes changed) by me!*
 - <https://brinckerhoff.org/clements/2194-csc431/essentials-of-compilation.pdf>
- Please stay *in the class* even if you think you are not prepared (talk to me if so)
 - Grading designed to be somewhat forgiving, projects done in groups (if you want only)



Why take CIS531?

- Compilers, like OS, involve nearly every aspect of computer science:
 - Algorithms (register allocation, program analysis, ...)
 - Data structures (syntax trees, control-flow graphs, ...)
 - Computability (grammars, automata, ...)
 - Programming languages (semantics, type checking, ...)
 - Systems (assembly code, ABIs, linking, ...)
 - Security (stack smashing, reverse engineering, ...)
 - Software Engineering (debugging, testing, specifications, ...)
- **We will talk (at least a little) about all of these!**

Why take CIS531? ...

- Great way to demonstrate to employers that you are capable of doing hard programming projects
- Distinguishing course that tests your knowledge of the rest of the curriculum
- Understand the internals of compiler design
 - Helps teach “how to learn” other languages
- Challenging (but very useful) exercise in debugging across layers
- Exposure to reading technical papers
- (Optional) learn how to work on teams to develop a large project

Grading

- There will be:
 - (50%) 4 projects — groups of up to three
 - (15%) Present a paper — I have selected papers, you will make slides
 - (30%) Exams — two written midterms
 - Not quite like traditional exams, we will talk about this
 - (5%) Participation — some incentive for you to come to class
 - Assessed via participation quizzes:
 - You will get 5% if you attend at least 2/3 of participation quizzes

Projects

- Project 1 (2 weeks) – Racket warmup, simple interpreter
- Project 2 (2 weeks) – Compiling arithmetic + variables to x86
- Project 3 (2 weeks) – Branching control-flow,
- Project 4 (~2.5 weeks) – set!, assignment conversion, procedures
- Project 5 – Final Project

Projects: Details

- You **can** work in groups of **up to three**
 - If you don't feel confident, *please team up with someone who does*
 - If you are more confident, please consider working with someone who is a beginner
 - There are some **special rules** (more work) required for teams
- Projects coded up in either Racket or Python (starter files for both)
 - I **strongly** recommend Racket, and I will teach it during class
 - Also, I will give a significant amount of hints, and will code up a significant portion of the solution in class for students to follow along.
 - Racket may look a bit intimidating, but it's an easy language—I have a set of course videos which will bring you quickly up to speed if you want to learn.
- Projects graded via **autograder**

Weekly Topics (first half of course...)

- **Week 1:** Introduction, preliminaries, x86 assembly crash course
- **Week 2:** Compiling basic blocks to straight-line assembly
- **Week 3:** Register allocation, liveness analysis, interference graphs
- **Week 4:** Register allocation via graph coloring
- **Week 5:** Branching control-flow (if/cond/...)
- **Week 6:** Loops and data-flow analysis
- **Week 7:** Program analysis, lattices, abstract interpretation
 - Midterm exam here

Weekly Topics (second half of course...)

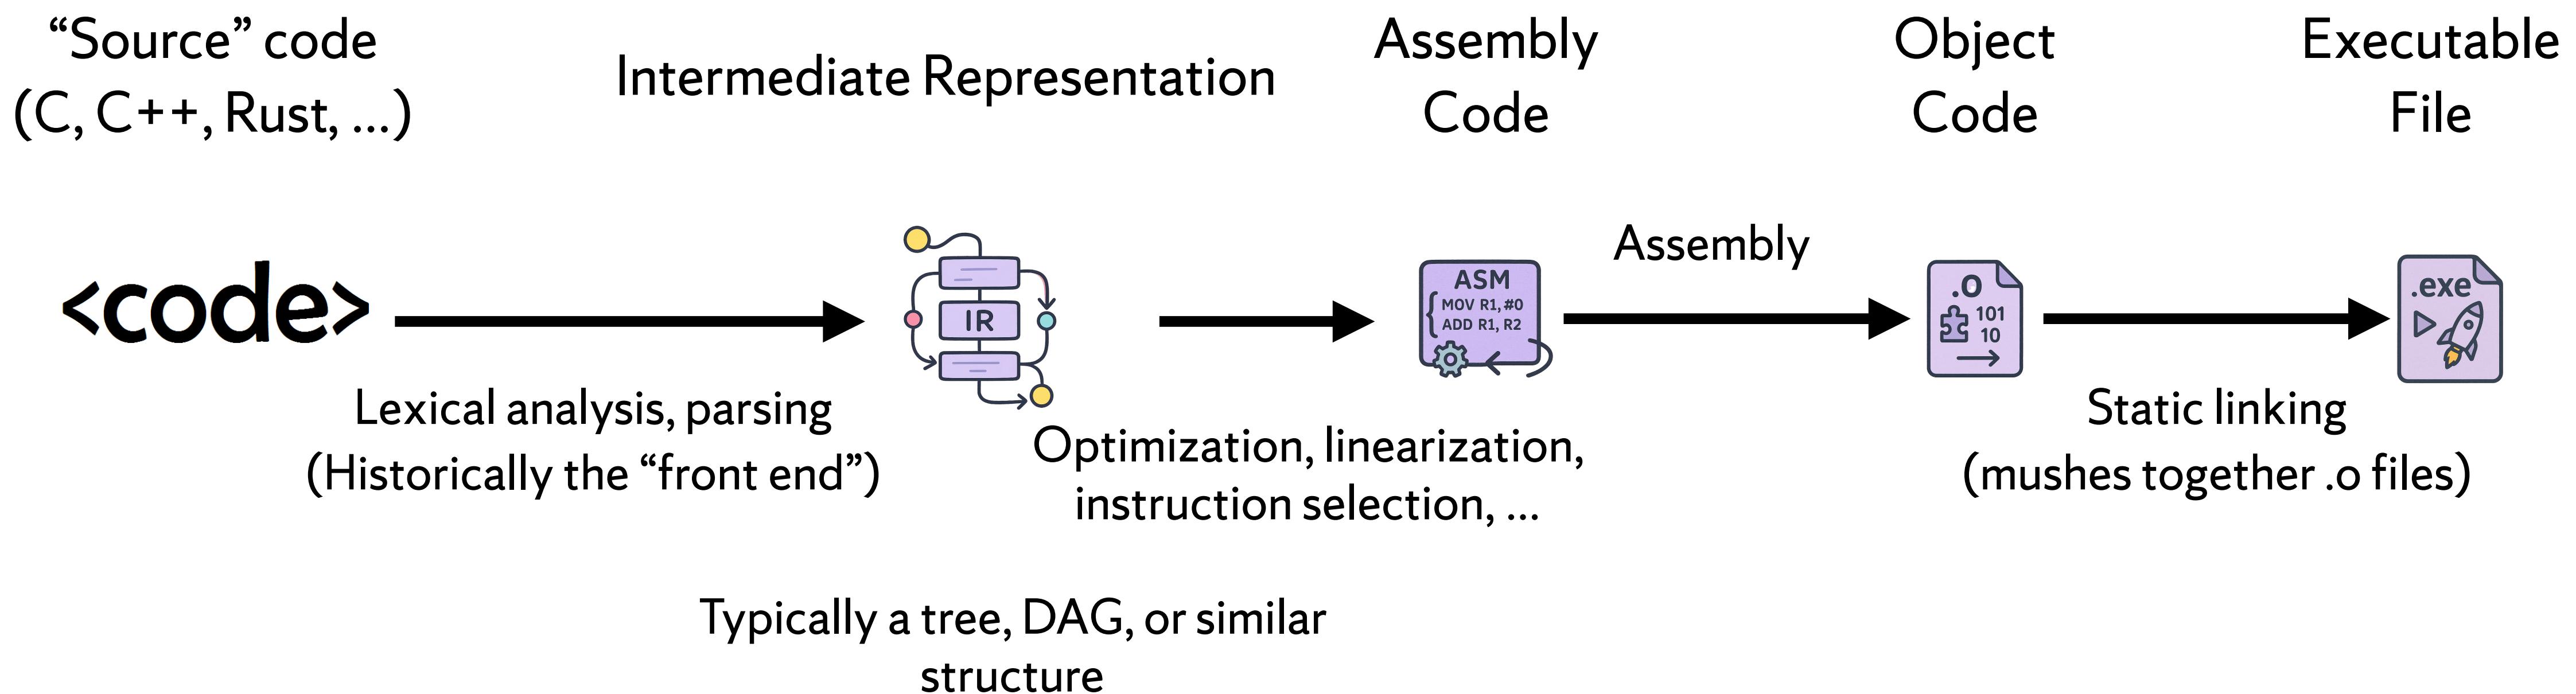
- **Week 8:** Tuples and Garbage Collection / Allocation
- **Week 9:** More GC, runtime systems
- **Week 10:** Functions, direct calls
- **Week 11:** Multiple parameters, tail calls, indirect calls
- **Week 12:** Advanced topics: Closure conversion
- **Week 13:** Advanced topics: Object orientation
- **Week 14:** Advanced topics: Security and Decompilation

Group Details

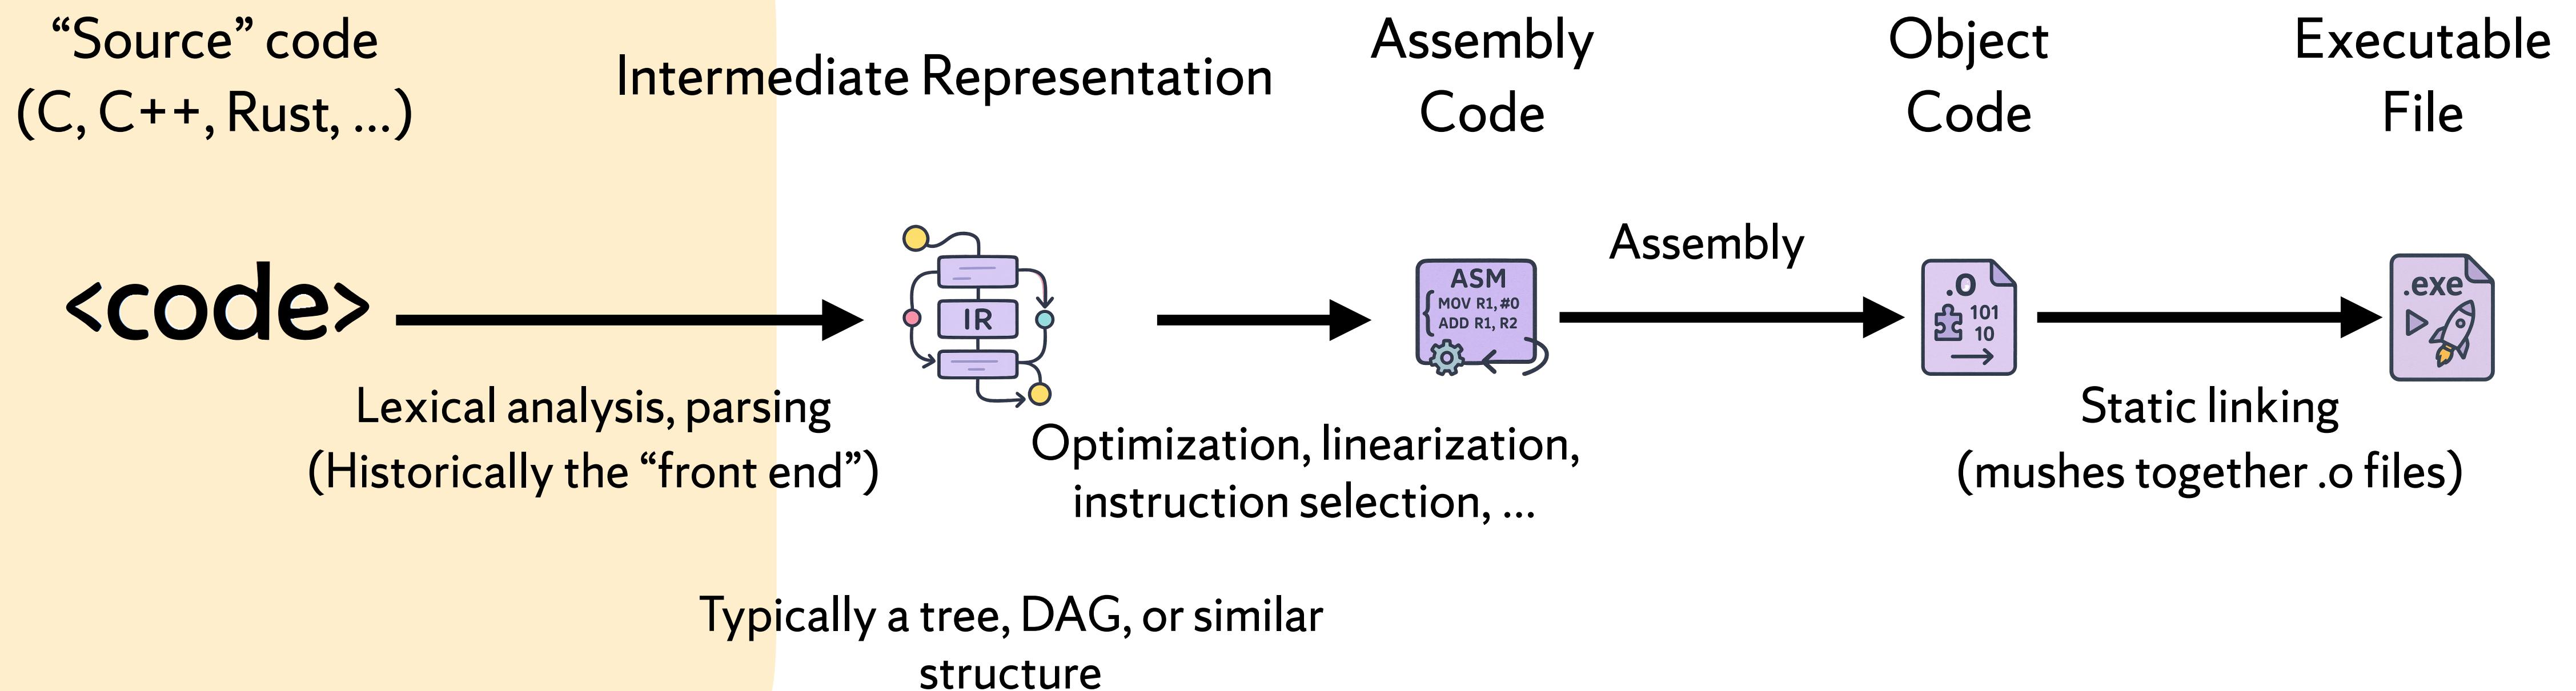
- Groups are per-project
- If you want to work in a group, you have to commit to meeting regularly
- You can work asynchronously, but you have to do the following:
 - You must spend **half** of your time writing / reviewing code synchronously, together
 - This can include debugging, or talking through concepts / white boarding
- For each group where you work as a group, each group member must submit an email to me at the end of the project saying:
 - What they personally contributed to the project?
 - Were they satisfied enough with their group's performance?
 - What things could they have done personally to improve group performance?
 - What could others in their group do to improve group performance?
 - After each project, the group must discuss: what are the top two things we can commit to doing to improve our collaboration on the *next* project (if we worked together again)?

What is a compiler?

Historically, compiler separated into “frontend” (lexing, parsing, etc...), “middle-end” (semantic analysis, IR construction, some optimization), and “backend” (dumping to machine-specific IR, assembly, linking)

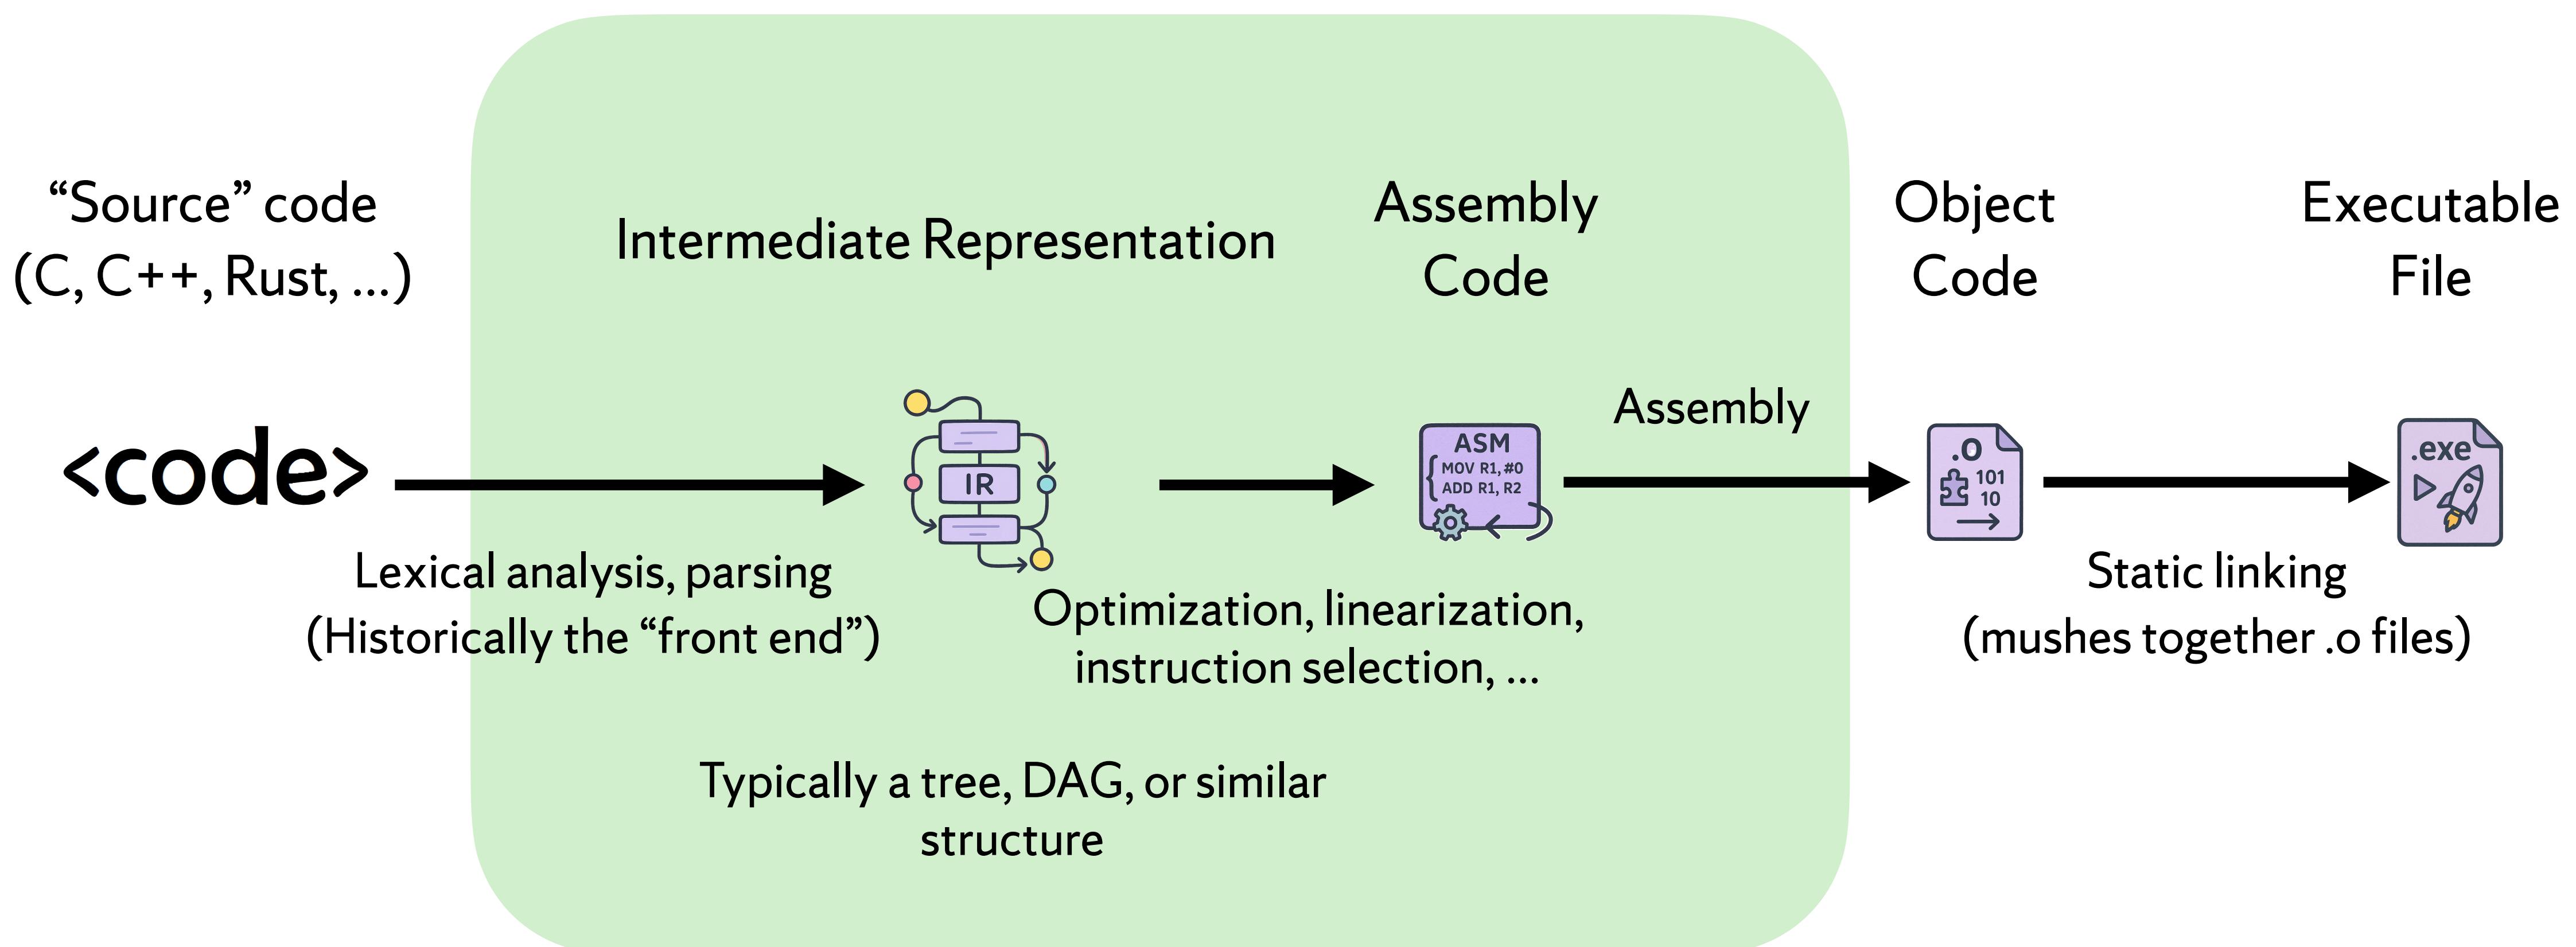


We will spend only a little time on the frontend, talking about parsing/lexing for only a few lectures



We will spend **almost all of our time** here. Going from the assembly code to object code is tedious and not very intellectually interesting: lots of assemblers / linkers exist already

Linking is interesting and hard, but we will not talk about it too much (not my expertise)



Compilers are designed in passes

- 
- Closeness to machine code
- Lexical analysis – separate out individual *tokens* from a stream
 - Syntactic analysis (parsing) – build AST / parse tree
 - Semantic analysis (name resolution, typing, etc...)
 - Intermediate Representation (IR) construction
 - Machine-independent optimization passes (possibly many) on IR
 - Machine-dependent lowering (addressing modes, vector ops, etc.)
 - Register allocation, spilling, etc.
 - Instruction scheduling, peephole optimization, etc.
 - Assembly / object code emission
 - Linking and relocation
 - Metadata emission (DWARF, ...) – GC metadata emission, debugging symbols, etc.

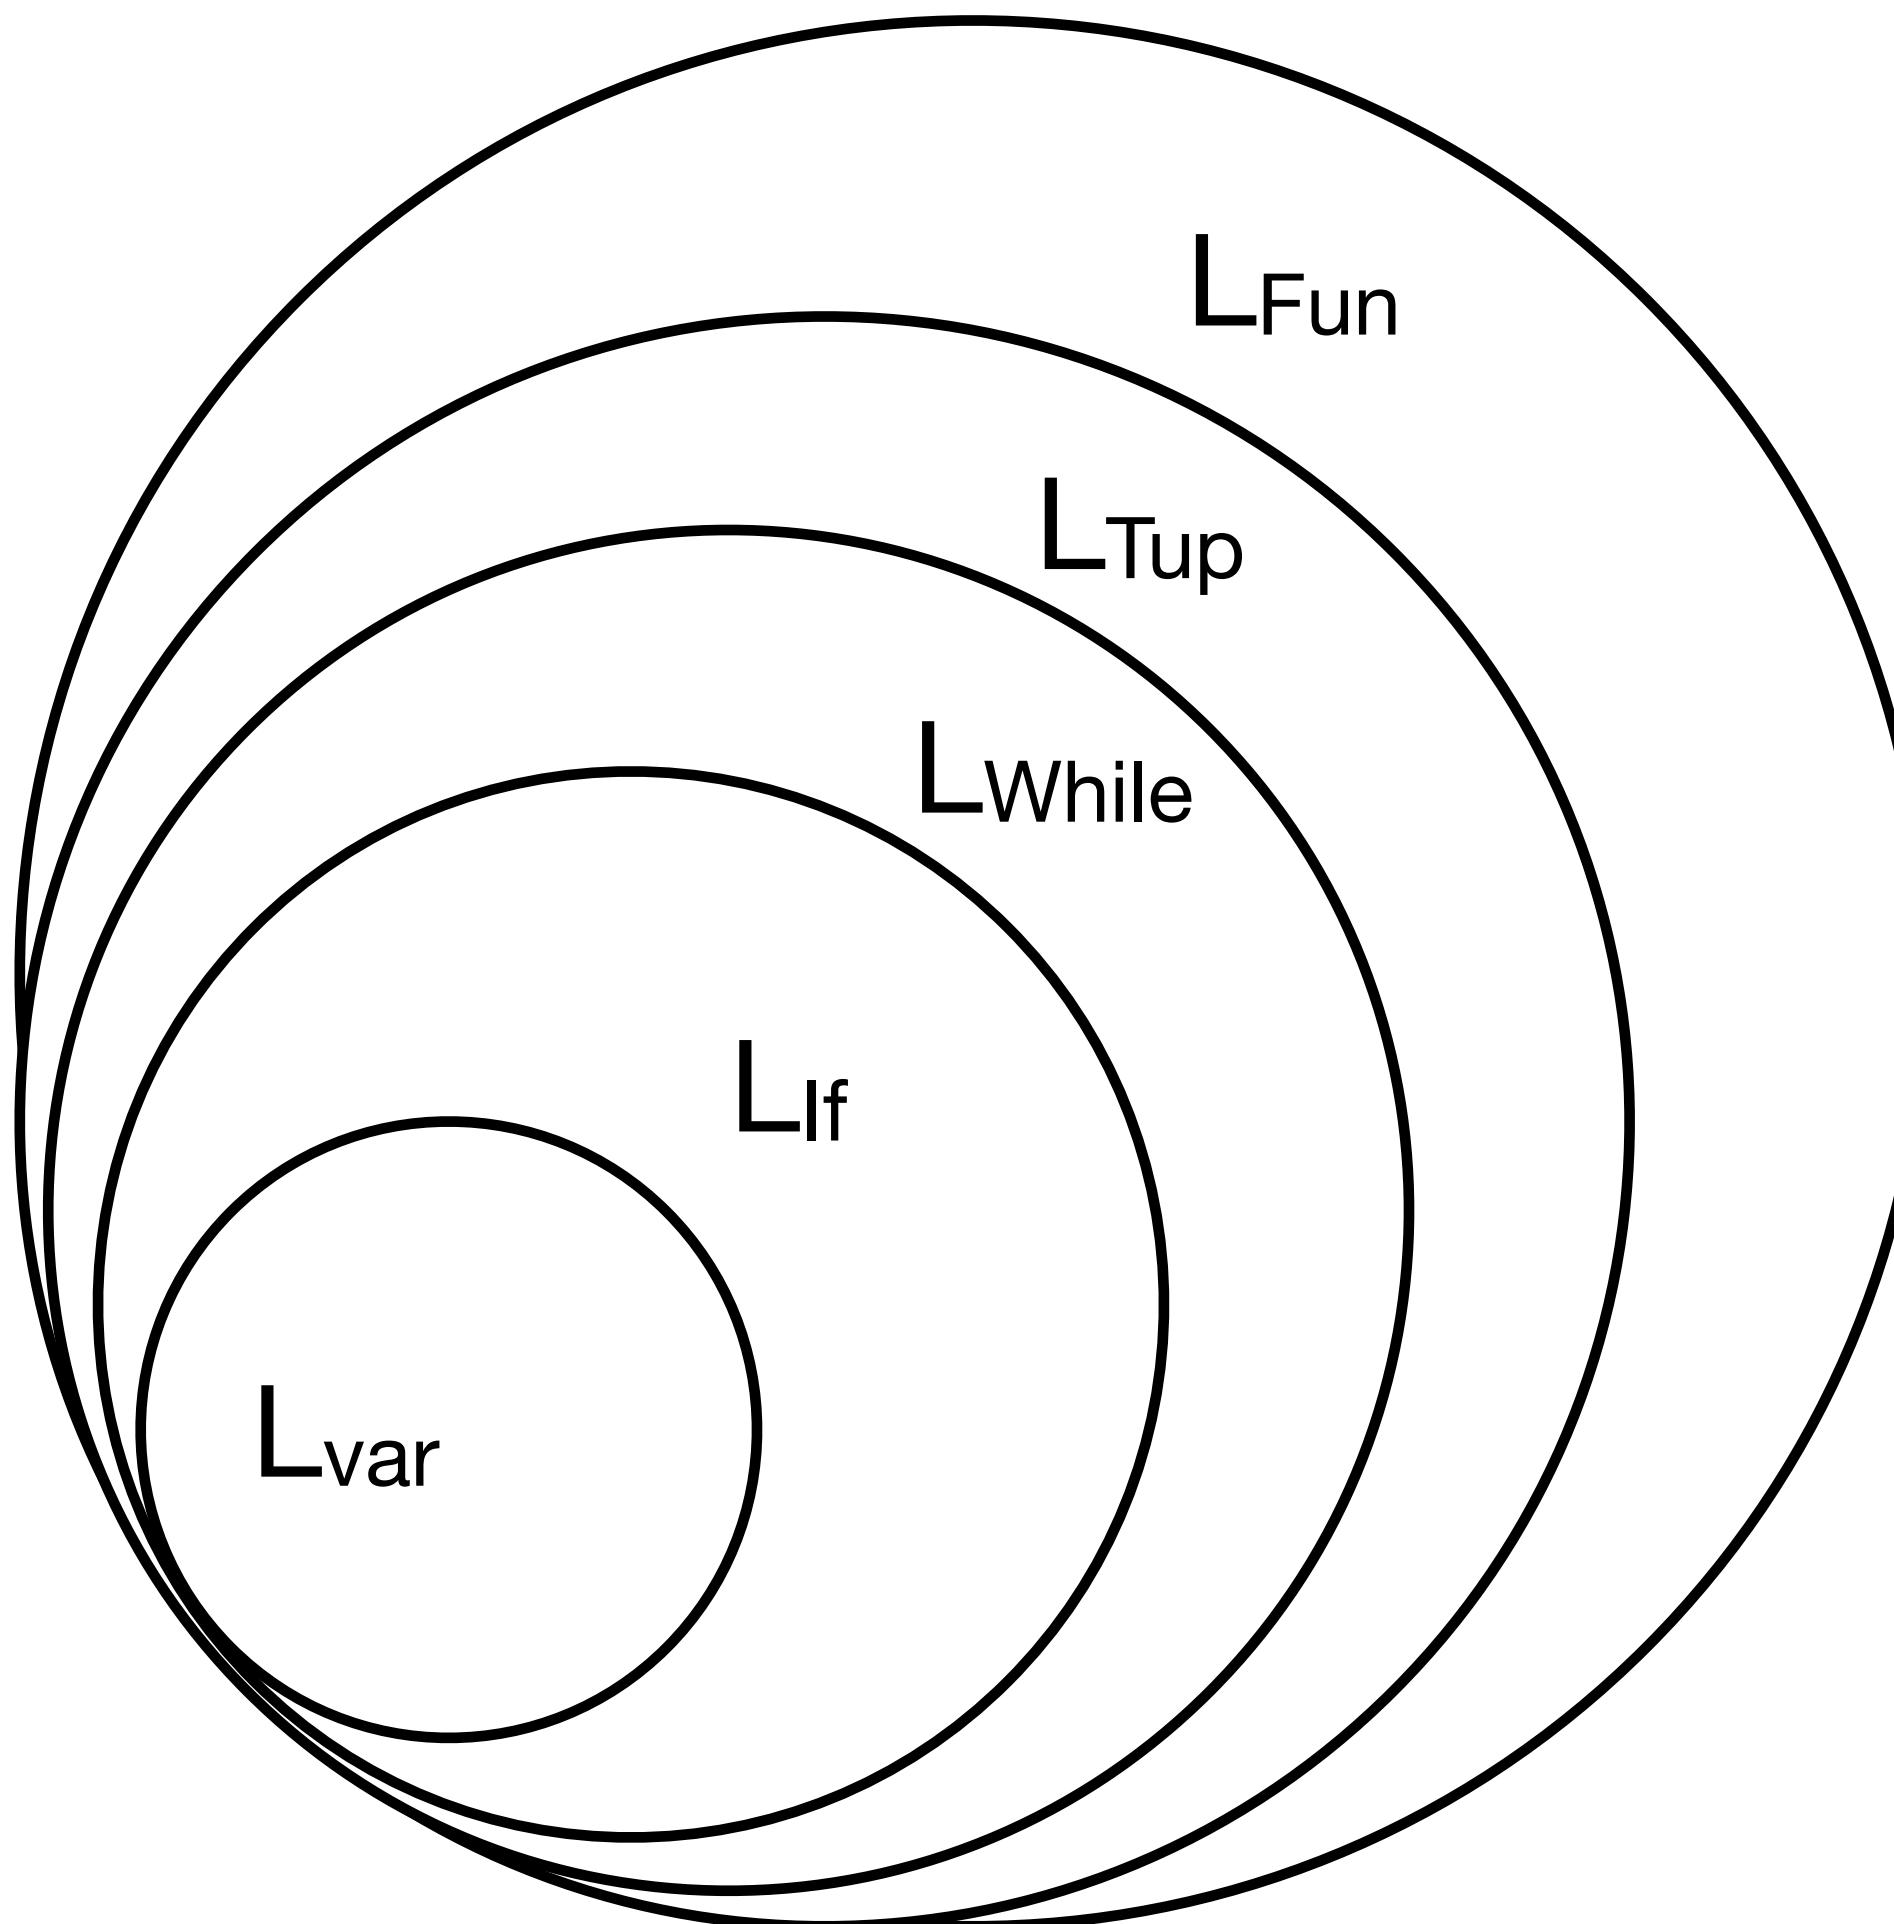
The Waterfall (Dragon Book) Approach: Pros/Cons

- Many classes take a top-down approach, but this has several drawbacks:
 - Lots of work, need to know whole language up-front, redoing lots of preexisting passes
 - Hard to beat modern C/C++/Rust compilers in practice (compile to C/C++/Rust?)
 - *The payoff is a long time down the road: can't run **any** program until the very end*
- Several possible solutions:
 - Write interpreters for intermediate IRs (**unsatisfying, seeing x86 is exciting!**)
 - Compile to higher-level language (C/C++/Rust)
 - Modern compilers (clang, etc.) are **amazing** in practice
 - Downside: you don't learn to truly compile to assembly, which is still relevant
 - But a very popular solution in practice (e.g., Soufflé Datalog)
 - Start in the middle (ignore parsing) by using an easily-parsable thing (e.g., JSON)
 - Parsing is interesting but orthogonal to the semantics

We take an *incremental* approach

- My course follows Jeremy Siek's approach (replicated at many other unis!)
- Unlike the traditional approach, we build a **whole** compiler in **every** project
- This way, you can use the compiler (completely) at every point
 - The **language** is the thing that grows, rather than the compiler
- Also, you **have** to develop a good debugging strategy (for assembly) early
 - It is crucial to acquire the skills necessary to debug the IR at multiple stages!

A nested tower of increasingly-powerful languages



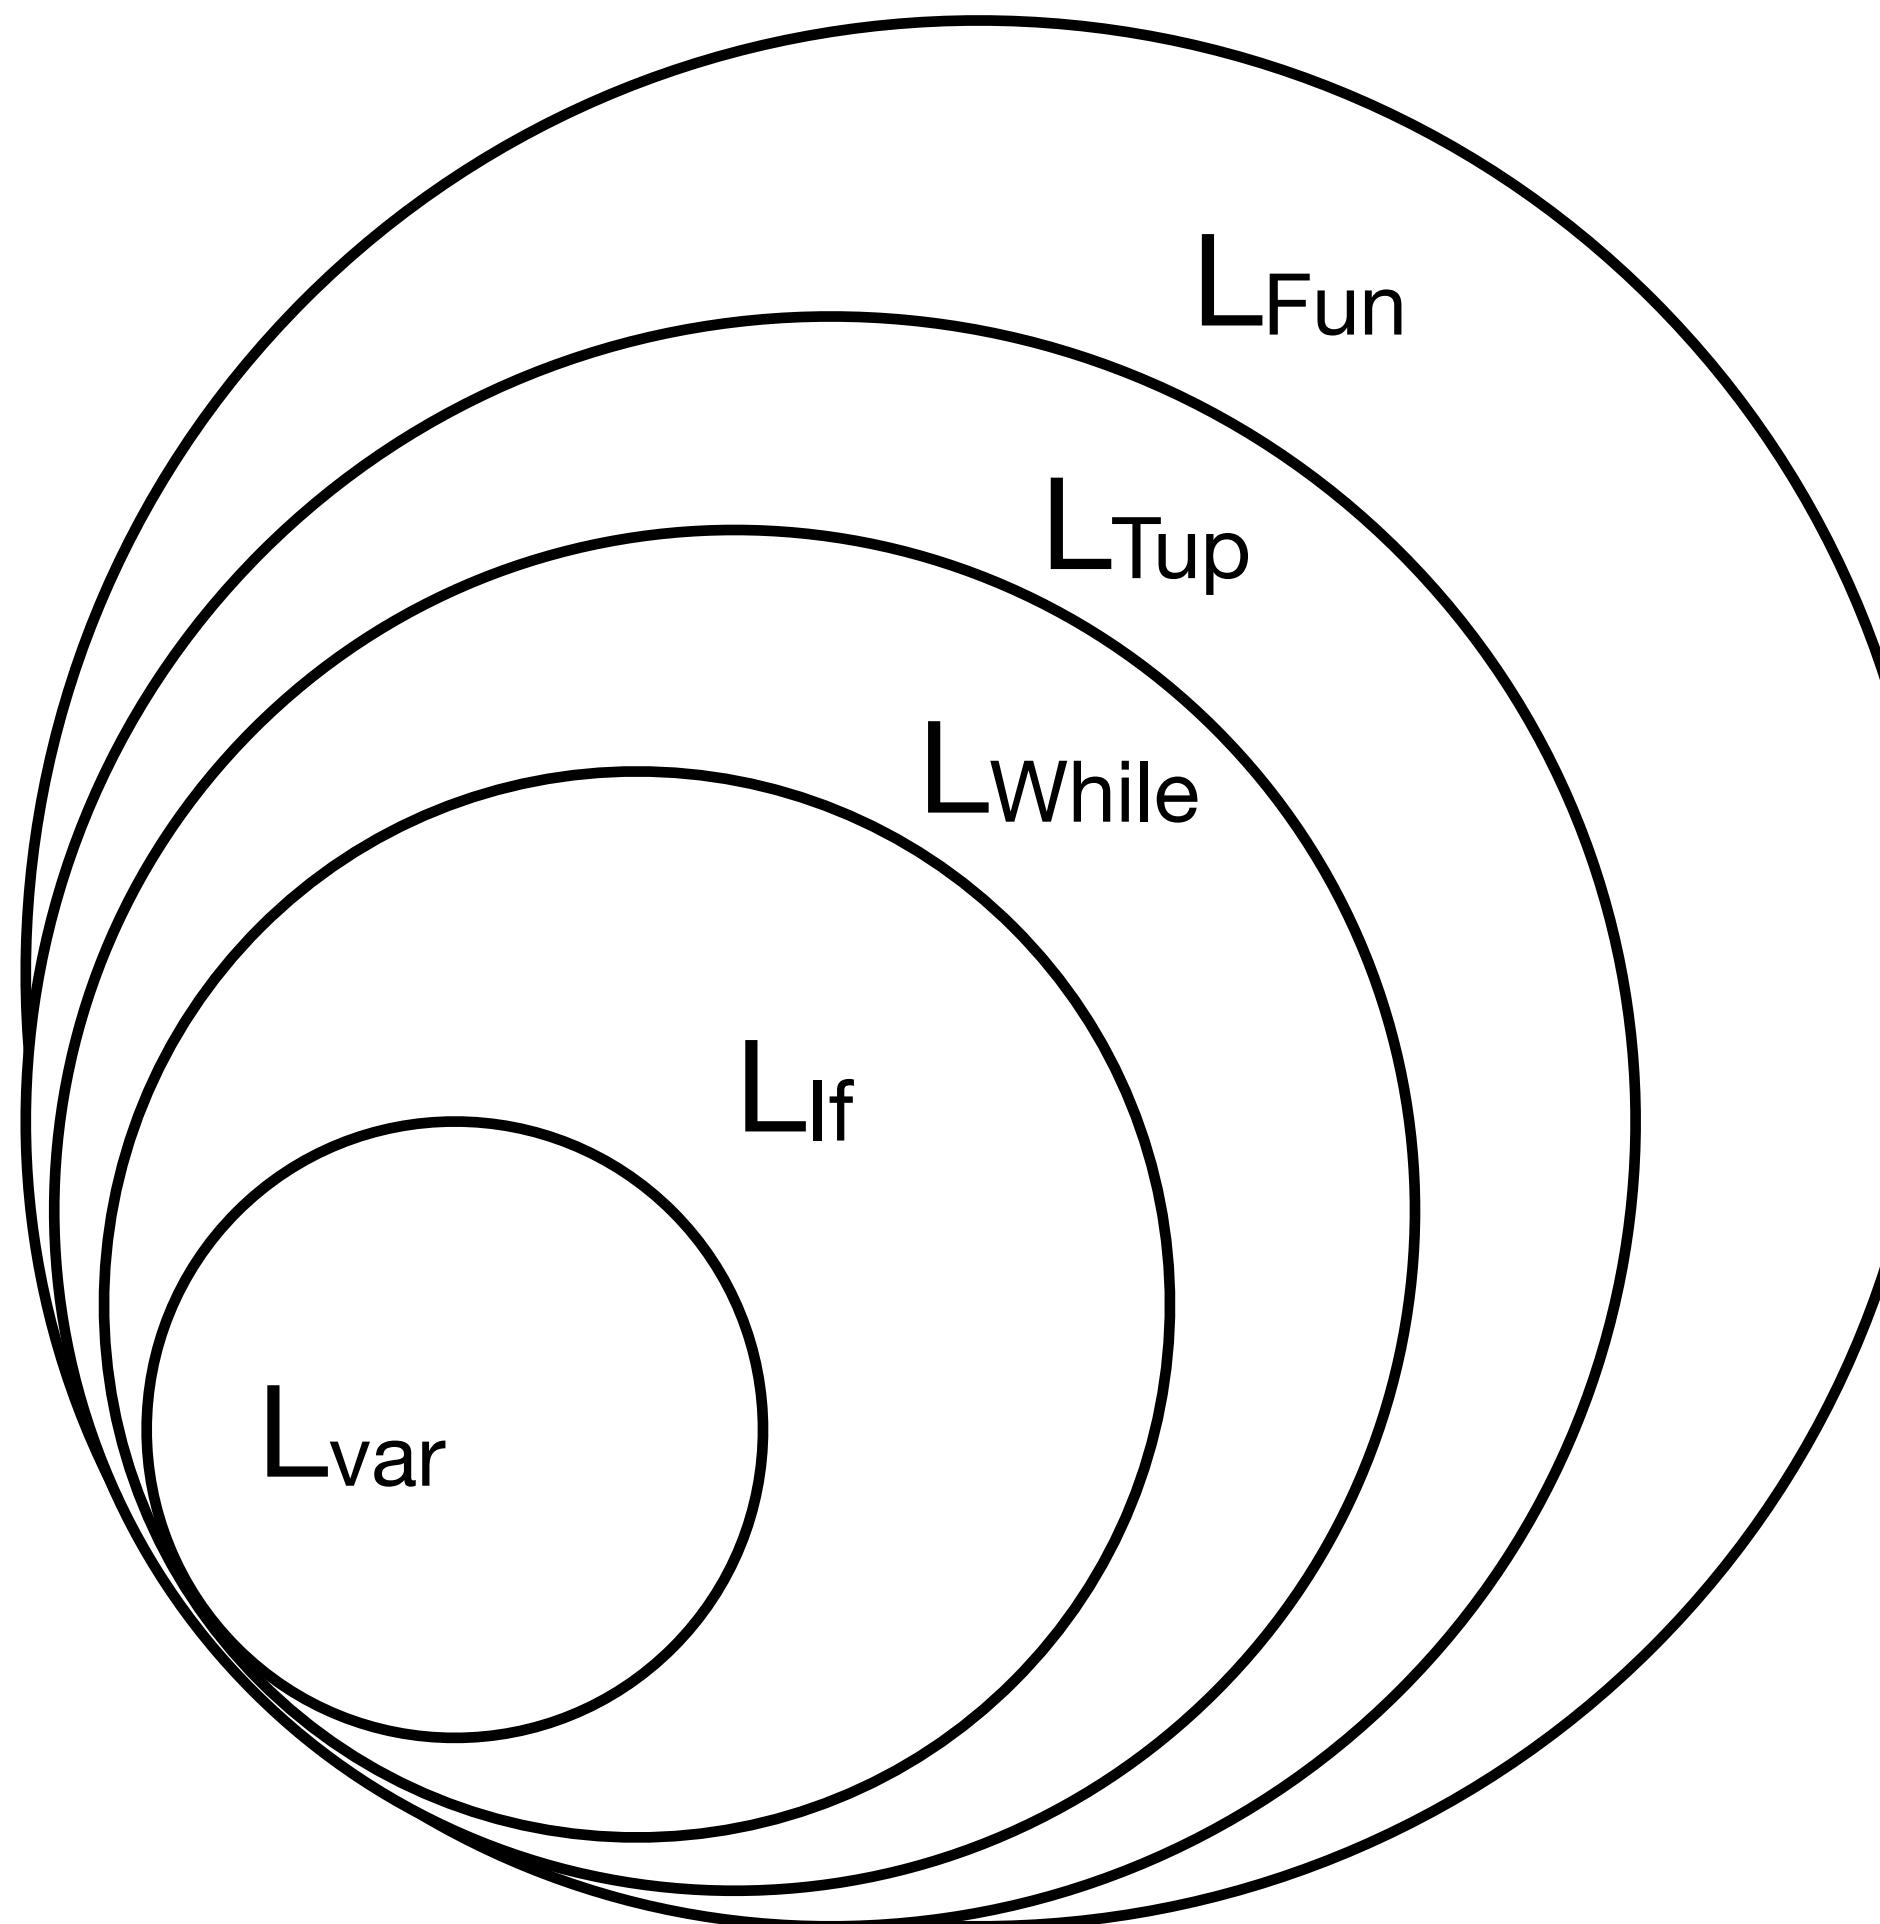
(Stolen from Jeremy Siek's slides)

Order of language the book presents:

- L_{Var} – integer +/-, straight-line variables
- L_{If} – branching control
- L_{While} – loops
- L_{Tup} – heap-allocated data, GC
- L_{Fun} – functions
- Some more...

We will follow roughly this structure, some variation toward the end (final project)

We have Racket (Scheme) and Python variants



Racket variant...

```
(let ([x (* 5 3)])
  (let ([y (+ x (* x x))])
    (+ (read) 5)))
```

Python variant...

```
x = 5 * 3
y = x + x*x
read_int() + 5
```

The key insight is that the **specific syntax doesn't really matter**.

We give Racket (Scheme) and Python variants, but only the surface syntax differs

We provide starter code and test harnesses in Python and Racket

I strongly recommend coding in Racket! I will offer substantial help (both in class / OH)!

Racket variant...

```
(let ([x (* 5 3)])
  (let ([y (+ x (* x x))])
    (+ (read) 5)))
```

Python variant...

```
x = 5 * 3
y = x + x*x
read_int() + 5
```

So what is the *language*?

- The focus of the course is **not functional programming!**
 - If you are concerned that you will only be learning useless FP stuff, **do not worry—the principles apply across languages / paradigms!**
 - We have test cases in Python **and** Racket (just syntactic differences!)
 - I am also open to you doing the class in other languages if you have a strong preference—though we **need to have a conversation beforehand** so I can grade your solutions
- I will program in Racket—not because it is a functional language—because it is the **best tool for the job**
 - *Do not worry* if you don't know Racket—I will keep it very simple, and we will go slow at first
 - *All of the topics in the course translate directly to Python, Rust, Java, C++, ...!*

Check your Knowledge

A few questions after this lecture:

- Is attendance required? If so, how is it measured?
- What language will be used for projects?
- Are we using a book? If so, do you have to pay for it?
- Can you work in groups? If so, what are the parameters?

Ask me (kkmicins@syr.edu) if you want to discuss answers after class